

Siddhesh Shankar

shankarsiddhesh1606@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Professional with 2+ years of experience turning complex datasets into measurable business outcomes across FinTech and E-Commerce. Proficient in SQL, Python, Power BI, Tableau, and ETL pipelines. Incoming MS Information Management student at UIUC (Fall 2026).

EDUCATION

University of Illinois, Urbana-Champaign (UIUC)

Champaign, IL

M.S. in Information Management (MSIM)

Aug. 2026 (*Incoming*)

Symbiosis Institute of Technology (SIT)

Pune, India

B.Tech in Computer Science & Engineering, Honors in Data Science

Apr. 2019 – Jun. 2023

EXPERIENCE

Business Intelligence Engineer

Sep. 2025 – Present

Synerpro Solutions LLP

- Built an analytics dashboard integrating website behavioral data to profile and classify traffic by industry, intent, and engagement signals, enabling the founding team to systematically shortlist and qualify high-potential clients.
- Implemented on-page SEO strategies including keyword structuring, meta tag optimization, and performance tuning, improving organic search rankings and driving measurable traffic growth from target business segments.
- Designed and developed the company website end-to-end using HTML, CSS, and Vanilla JS, and established a LinkedIn and blog-based content strategy to support early-stage client acquisition.

Data Analyst

Jul. 2023 – Aug. 2025

Oracle Financial Services Software Ltd. (OFSS)

- Analyzed 70,000+ daily financial transactions (MT940/MT950/MT942) in a Nostro Reconciliation System, improving transaction accuracy by 25% through root-cause analysis and process optimization.
- Engineered real-time Nostro Reconciliation dashboards for 150–200 internal ops and reconciliation users at HDFC Bank (Corporate), implementing core matching logic in PL/SQL on OracleDB and leveraging Apache Kafka and Java Message Service (JMS) for event streaming, tracking 3,500+ TPS and eliminating manual LOOKUP-based cross-checking of SWIFT MT942 against SWIFT MT950/MT940 records.
- Streamlined data flow between legacy SOAP systems and RESTful services, reducing data integration turnaround time by 40% and improving cross-system data consistency for downstream reporting.

Business Analyst

Jan. 2023 – Jun. 2023

EasyEcom Pvt. Ltd.

- Optimized SQL queries and automated ETL pipelines in a high-volume e-commerce environment, cutting manual data processing time by 30% and improving reporting accuracy for business-critical dashboards.
- Built automated data-fetching workflows using Google Apps Script, reducing dashboard refresh time by 40% and enabling daily real-time reporting for supply chain and growth teams.
- Instrumented end-to-end user behavior tracking via Amplitude for D2C clients including Meesho, Flipkart, and ManMatters, translating funnel and retention data into product decisions that shaped feature prioritization.

PROJECTS

QueryMind | *Python, LangChain, SQL, NoSQL, Pandas*

Mar. 2026 – Present

- Developing a solo Python CLI tool that accepts natural-language questions over CSV and multi-sheet Excel files and returns structured analytical insights, removing the need for ad-hoc scripting during exploratory analysis.
- Built a hybrid rule-based and LLM intent interpreter that maps analyst queries (top-N, comparison, trend, aggregation) to structured execution plans using pandas, handling edge cases such as packed integer dates, BOM-encoded files, and cross-sheet merges.

Autonomous Data-Analyst | *Python, FastAPI, Streamlit, LangChain*

Feb. 2026 – Present

- Engineered an agentic AI pipeline using multi-agent architecture (Cleaning, EDA, Insight, Report agents) to automate end-to-end data analysis, reducing manual analysis effort by 70%.
- Developed a FastAPI backend for CSV/Excel ingestion generating structured insights in under 5 seconds, with a Streamlit UI producing downloadable McKinsey-style PDF reports with dynamic charts.

InfoGuard-AI | *Python, MongoDB, GitHub Actions*

Jan. 2026 – Mar. 2026

- Engineered a real-time AI monitoring pipeline processing 100+ Wikipedia pages/cycle at 90+ pages/min, detecting semantic anomalies via NLP risk scoring (MiniLM embeddings) across 300+ production runs.
- Designed a multi-factor risk engine combining semantic similarity, content signals, and user behavior, reducing false positives and maintaining consistent performance under 16 min average runtime.

TECHNICAL SKILLS

Languages & Querying: Python, R, SQL (PostgreSQL, OracleDB), NoSQL (MongoDB)

Analytics & BI: Power BI, Tableau, Looker Studio, MS Excel (Pivots, XLOOKUP), Amplitude, Google Analytics

Libraries & Frameworks: pandas, NumPy, Matplotlib, Seaborn, FastAPI, Streamlit, LangChain

Data Engineering: ETL Pipelines, Google Apps Script, Data Cleaning & Transformation

Tools: Git, Docker, Jupyter Notebook, GitHub Actions, VS Code